

Knowledge Graph Project (FORWARD Data Lab)



Part 1. ETL Pipeline

Load libraries

```
In [1]: import pyspark as sp
import numpy as np
import string
import re
```

Set environment variable (optional)

```
In [2]: import os
os.environ["PYSPARK_PYTHON"]="/usr/bin/python3"
os.environ["PYSPARK_DRIVER_PYTHON"]="/usr/bin/python3"
```

Setup cluster

```
In [3]: from pyspark.sql import SparkSession
from pyspark import SparkContext

spark = SparkSession \
    .builder \
    .appName("Python Knowledge Graph Project") \
    .config("spark.driver.memory", "55g") \
    .master("local[*]") \
    .getOrCreate()
```

Load dataset

- Documents: 1.7m documents (arXiv dataset)
- Keyword list: 100,000 keywords (FORWARD Data Lab)
- Download your documents here: <https://www.kaggle.com/Cornell-University/arxiv>.

```
In [4]: df_keywords = spark.read.csv("./mag_cs_keywords.csv", header=True)
df_arxiv = spark.read.json("./arxiv-metadata-oai-snapshot.json")
```

```
In [5]: df_keywords = df_keywords.repartition(16)
df_arxiv = df_arxiv.repartition(16)
```

```
In [6]: df_keywords.rdd.getNumPartitions()
```

```
Out[6]: 16
```

```
In [7]: df_keywords.count()
```

```
Out[7]: 104654
```

```
In [8]: abstracts = df_arxiv.select("id","abstract")
keywords = df_keywords.select("normalizedName")
```

Cleaning & Normalization

```
In [9]: from pyspark.sql.functions import udf, col
from pyspark.sql import Row
from pyspark.sql.types import ArrayType, StructField, StructType, StringType,

# get only lowercase alphabets
def strip_non_ascii(data_str):
    ''' Returns the string without non ASCII characters'''
    stripped = (c.lower() for c in data_str if 96 < ord(c.lower()) < 123 or ord(c.lower()) > 123)
    return ''.join(stripped)
# setup pyspark udf function
strip_non_ascii_udf = udf(strip_non_ascii, StringType())
```

```
In [10]: abstracts = abstracts.withColumn('normalized', strip_non_ascii_udf(abstracts['abstract']))
abstracts = abstracts.select("id","normalized")
```

```
In [11]: abstracts.show()
```

id	normalized
0705.2030	the group eso s...
0801.0970	this paper is c...
0709.3236	we present a ge...
0809.3030	the tragedy of ...
0712.0290	the ability to ...
0805.2349	this article pr...
0805.1292	we turn the chu...
0712.3567	the random sequ...
0705.1654	a smallscale tr...
0803.3867	we present a ch...
0709.0351	for precision s...
0711.4659	in a newly intr...
0705.1358	we propose two ...
0710.2002	we present prec...
0801.1957	we calculate th...
0707.0756	many topologica...
0808.0066	the reasonablen...
0801.3027	in this paper w...
0707.0344	in this paper w...
0807.4946	we establish lo...

only showing top 20 rows

Keyword List for Query

- `keywords_dict` : Hash table is used to find whether documents contain keywords
- My first attempt was using regular expression (`re.findall()`).
- When using regular expression, there could be a time complexity issue due to unnecessary computational overhead
- To prevent optimize algorithm, we should search `from document to word`.

```
In [12]: keywords_list = list(keywords.select('normalizedName').toPandas()['normalized'])
keyword_dict = {word:0 for word in keywords_list}
```

Extract Keyword from the keyword list

```
In [13]: import re, nltk

# Define the function you want to return
def extract(s):
    all_matches = set()

    ngrams = list()
    ngrams.append(s.split(" "))
    for i in range(1,12): # all ngrams in keyword list
        ngrams.append(nltk.ngrams(ngrams[0], i))

    for ngram in ngrams:
        for token in ngram:
            token = ' '.join(token)
            if keyword_dict.get(token, -1) >= 0:
                keyword_dict[token] += 1
                all_matches.add(token)
    return ','.join(all_matches)

# Create the UDF, note that you need to declare the return schema matching the
extract_udf = udf(extract, StringType())

# Apply it
df = abstracts.withColumn('extracted', extract_udf(abstracts['normalized']))
```

```
In [14]: df.show()
```

id	normalized	extracted
0705.2030	the group eso s...	power,momentum,ro...
0801.0970	this paper is c...	selection,minimax...
0709.3236	we present a ge...	perturbation theo...
0809.3030	the tragedy of ...	set,crowdsourcing...
0712.0290	the ability to ...	mechanism,morphol...
0805.2349	this article pr...	liquid helium,pha...
0805.1292	we turn the chu...	physics,computati...
0712.3567	the random sequ...	deposition,initia...
0705.1654	a smallscale tr...	thermal,temperatu...
0803.3867	we present a ch...	product,term
0709.0351	for precision s...	field,perturbatio...
0711.4659	in a newly intr...	square,phase,posi...
0705.1358	we propose two ...	phase,phase trans...
0710.2002	we present prec...	pair production,d...
0801.1957	we calculate th...	quantum,phase,fie...
0707.0756	many topologica...	generalization,se...
0808.0066	the reasonablen...	perturbation theo...
0801.3027	in this paper w...	brownian motion,p...
0707.0344	in this paper w...	bridge,is a,devia...
0807.4946	we establish lo...	inflow,outflow,la...

only showing top 20 rows

$$PMI(A, B) = \frac{P(A, B)}{P(A) * P(B)}$$

$$P(A, B) = \frac{c((A, B), C_{\text{pairs}})}{c(C_{\text{pairs}})}$$

$$P(A) = \frac{c(A)}{c(C)}, P(B) = \frac{c(B)}{c(C)}$$

where C stands for word collection

1. Get Word Count

- `corpus` : It is a dictionary of keywords

```
In [15]: corpus_temp = df.select("extracted").rdd.repartition(16).flatMap(lambda x: (x
corpus = corpus_temp.countByValue()
```

```
In [53]: len(corpus)
```

```
Out[53]: 70342
```

2. Get Cooccurrence Count

- `corpus2` : It is a dictionary where key is tuple (pair of keywords) and value is frequency

```
In [16]: rdds = df.select("extracted").rdd.repartition(16).map(lambda x: x['extracted'])
```

```
In [17]: from itertools import combinations
corpus2_temp = rdds.flatMap(lambda x: combinations(x, 2))
corpus2 = corpus2_temp.countByValue()
```

```
In [54]: len(corpus2)
```

```
Out[54]: 15617821
```

3. How to avoid bottleneck?

- Use clusters! (Amazon EMR, S3)

4. Save Result Here

- Uncomment to save to text file

```
In [18]: # corpus_temp.repartition(16).saveAsTextFile("./corpus_ngram.txt")
# corpus2_temp.repartition(16).saveAsTextFile("./corpus2_ngram.txt")
```

5. Load From File

- Uncomment to load from file

```
In [ ]: # corpus = spark.sparkContext.textFile("./corpus_ngram.txt").countByValue()
# corpus2 = spark.sparkContext.textFile("./corpus2_ngram.txt").countByValue()

"""
Clean corpus when read from file
Since we are loading from a textfile, tuples are converted into strings
"""

# from collections import defaultdict
# corpus_temp = defaultdict(int)

# corpus = dict(map(lambda x: (x[0].strip(), x[1]), corpus.items()))
# for k, v in corpus2.items():
#     k1, k2 = re.sub("'", "", k.strip('"()')).split(",")
#     k2 = k2.strip()
#     w1, w2 = min(k1, k2), max(k1, k2)
#     corpus_temp[(w1, w2)] += v
# corpus2 = corpus_temp
```

Part 2. Graph Visualization

Draw Graph

```
In [19]: freq_threshold = 0
PMI_threshold = 10
```

```
In [20]: import networkx as nx
import math
import pdb
G = nx.Graph()

total_word_freq = sum(corpus.values())
total_pair_freq = sum(corpus2.values())
for (n1, n2), cooccurrence in corpus2.items():
    if corpus[n1] > freq_threshold and corpus[n2] > freq_threshold and cooccurrence > PMI_threshold:
        pa = corpus[n1] / total_word_freq
        pb = corpus[n2] / total_word_freq
        pab = cooccurrence / total_pair_freq

        pmi = math.log(pab/(pa*pb), 2)
        G.add_edge(n1, n2, weight=pmi)
```

RUN QUERY HERE

- Type your query here
- Sample queries include
 - gene, entitiy relationship model, business concept, data mining

```
In [40]: query = 'entity relationship model'
sorted(G.adj[query].items(), key=lambda x: x[1]['weight'], reverse=True)[:5]
```

```
Out[40]: [('requirement document', {'weight': 21.543808138427178}),
          ('data entity', {'weight': 21.543808138427178}),
          ('database design', {'weight': 16.589611828040304}),
          ('business process', {'weight': 13.857307611243959}),
          ('business', {'weight': 10.290551558646024})]
```

```
In [41]: query = 'second language acquisition'
sorted(G.adj[query].items(), key=lambda x: x[1]['weight'], reverse=True)[:5]
```

```
Out[41]: [('second language writing', {'weight': 18.958845637706023}),
          ('contrastive analysis', {'weight': 18.958845637706023}),
          ('english as second language', {'weight': 17.37388313698487}),
          ('question analysis', {'weight': 16.958845637706023}),
          ('english as a second language', {'weight': 16.958845637706023})]
```

```
In [42]: query = 'bernoulli distribution'
sorted(G.adj[query].items(), key=lambda x: x[1]['weight'], reverse=True)[:5]
```

```
Out[42]: [('multiple outcome', {'weight': 13.698318087482804}),
          ('dynamic decision making', {'weight': 13.698318087482804}),
          ('multivariate gaussian model', {'weight': 13.43528368164901}),
          ('disorder problem', {'weight': 13.212891260312562}),
          ('simultaneous perturbation stochastic approximation',
           {'weight': 13.020246182370165})]
```

```
In [43]: query = 'business concept'
sorted(G.adj[query].items(), key=lambda x: x[1]['weight'], reverse=True)[:5]
```

```
Out[43]: [('customer segment', {'weight': 21.543808138427178}),
          ('strategic change', {'weight': 21.543808138427178}),
          ('research knowledge', {'weight': 18.736453216369576}),
          ('secondary data', {'weight': 16.685827143299605}),
          ('business model', {'weight': 15.168768707080254})]
```

```
In [52]: query = 'algorithm'
sorted(G.adj[query].items(), key=lambda x: x[1]['weight'], reverse=True)[:5]
```

```
Out[52]: [('attribute oriented induction', {'weight': 5.063633283582013}),
          ('coding algorithm', {'weight': 5.063633283582013}),
          ('resource allocation algorithm', {'weight': 5.063633283582013}),
          ('power control algorithm', {'weight': 5.063633283582013}),
          ('extraction algorithm', {'weight': 5.063633283582013})]
```

- Higher than 10 pmi demonstrates a strong relationship between the two words
- Below is the sample pmi table from the wikipedia dataset.

word 1	word 2	count word 1	count word2	co-occurence	pmi
puerto	rico	1938	1311	1159	10.0349081703
car	driver	5578	2749	1384	8.41470768304
it	the	283891	3293296	3347	-1.72037278119
of	and	1761436	1375396	1190	-3.70663100173

- Reference: https://en.wikipedia.org/wiki/Pointwise_mutual_information

Top 20 Related Words

- Queries and their related keywords are sorted here by pmi.
- Natural science domain seems to working well

```
In [22]: result = list()
for i, q in enumerate(list(G.nodes)):
    weight = sorted(G.adj[q].items(), key=lambda x: x[1]['weight'], reverse=True)

    if len(weight)==0:
        r1, r2, r3 = 0, 0, 0
        pmi = 0
    elif len(weight)==1:
        r1, r2, r3 = weight[0][0], 0, 0
        pmi = weight[0][1]['weight']
    elif len(weight)==2:
        r1, r2, r3 = weight[0][0], weight[1][0], 0
        pmi = weight[0][1]['weight']
    else:
        r1, r2, r3 = weight[0][0], weight[1][0], weight[2][0]
        pmi = weight[0][1]['weight']

    result.append((q, r1, r2, r3, pmi))

result.sort(key=lambda x: x[4], reverse=True)

for q, r1, r2, r3, _ in result[:20]:
    print("query: {}".format(q))
    print("result: {0}, {1}, {2}\n".format(r1, r2, r3))
```


query: gibberellin
result: brassinosteroid, principal mechanism, signalling pathways

query: brassinosteroid
result: gibberellin, principal mechanism, signalling pathways

query: customer segment
result: strategic change, business concept, research knowledge

query: strategic change
result: customer segment, business concept, research knowledge

query: business concept
result: customer segment, strategic change, research knowledge

query: digital single market
result: language grid, single market, multilingualism

query: language grid
result: digital single market, single market, multilingualism

query: poor nutrition
result: mental development, chatbot, immunity

query: mental development
result: poor nutrition, chatbot, immunity

query: coronary ligation
result: isoprenaline, ligation, myocardial infarction

query: isoprenaline
result: coronary ligation, ligation, myocardial infarction

query: synthetic drugs
result: lipolysis, thermogenesis, high morbidity

query: lipolysis
result: synthetic drugs, thermogenesis, high morbidity

query: requirement document
result: data entity, entity relationship model, database design

query: data entity
result: requirement document, entity relationship model, database design

query: entity relationship model
result: requirement document, data entity, database design

query: genetic linkage map
result: superphylum, polyphemus, ecdysozoa

query: superphylum
result: genetic linkage map, polyphemus, ecdysozoa

query: polyphemus
result: genetic linkage map, superphylum, ecdysozoa

query: conductive hearing loss
result: sound stimulation, umbo, eardrum

Visualization

- The final plot shows that related words form a group of cluster

```
In [23]: samples = ["spastic", "hypertonia", "reflex", "muscle", "endophyte", "jatroph
```

```
In [24]: matrix = list()
for i, row in enumerate(samples):
    row_temp = list()
    for j, col in enumerate(samples):
        w1, w2 = samples[i], samples[j]
        if G.adj.get(w1, 0) and G.adj[w1].get(w2, 0) and G.adj[w1][w2].get("w
            row_temp.append(G.adj[w1][w2]["weight"])
        else:
            row_temp.append(0)
    matrix.append(row_temp)
```

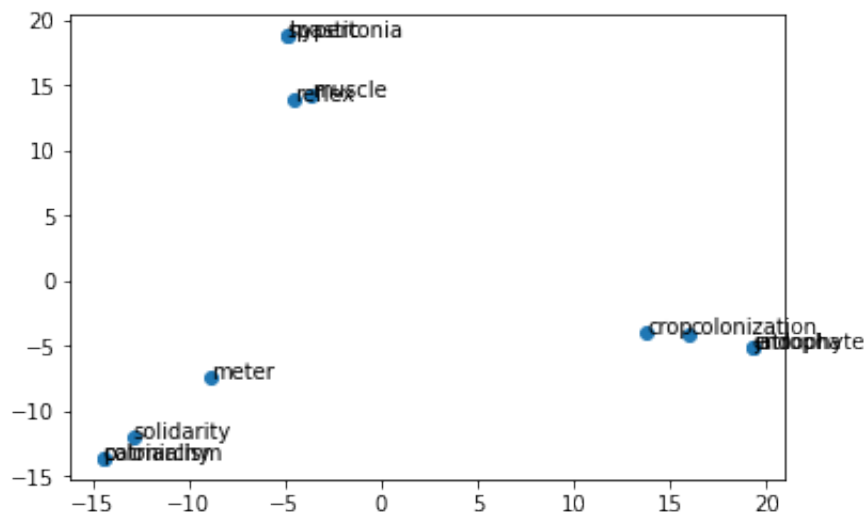
```
In [25]: from sklearn.decomposition import PCA
```

```
pca = PCA(n_components=2)
result = pca.fit_transform(matrix)
```

```
In [26]: from matplotlib import pyplot
```

```
pyplot.scatter(result[:, 0], result[:, 1])
for i, word in enumerate(samples):
    pyplot.annotate(word, xy=(result[i, 0], result[i, 1]))

pyplot.figure(figsize=(15,15))
pyplot.show()
```



<Figure size 1080x1080 with 0 Axes>

Future Directions

- Term Frequency
 - Apply TF transformation when counting words. (Sublinear tranformation , BM25)
 - Apply smoothing
- Embedding
 - Use cooccurence matrix to train word embeddings (Glove)
 - Or, load word embedding from Glove or Word2Vec and proejct to 2D plot.
 - Evalutate the similarity of two words by using l1 distance , l2 distance , or cosine similiarity .
- Different Approach
 - Try query expansion
 - Try information retrieval methods to infinitely expand the knowledge graph
- Metric
 - We may investigate NPMI
 - Entropy

References

- running pyspark in jupyter
 - <https://medium.com/@Vatsal410/running-pyspark-on-an-aws-cluster-through-jupyter-notebook-fe28f6ef47a4>
 - <https://medium.com/@christo.lagali/run-jupyter-notebooks-with-pyspark-on-an-emr-cluster-9630ef54c4e1>
 - <https://towardsdatascience.com/getting-started-with-pyspark-on-amazon-emr-c85154b6b921>
- reading and writing in Amazon s3
 - <https://docs.cloudera.com/runtime/7.0.2/developing-spark-applications/topics/spark-examples-of-accessing-s3-data-from-spark.html>
- textmining in apache spark
 - <https://runawayhorse001.github.io/LearningApacheSpark/textmining.html>
 - <https://medium.com/@serkansakinmaz/how-to-connect-amazon-s3-via-emr-based-pyspark-42707d540881>
- How to setup spark cluster in AWS
 - <https://aws.amazon.com/blogs/big-data/best-practices-for-successfully-managing-memory-for-apache-spark-applications-on-amazon-emr/>
- Extract keywords in document
 - <https://stackoverflow.com/questions/48869922/how-to-efficiently-check-if-a-list-of-words-is-contained-in-a-spark-dataframe>
 - <https://stackoverflow.com/questions/46410887/pyspark-string-matching-to-create-new-column>
- Useful tutorials
 - <https://spark.apache.org/docs/1.6.3/ml-features.html>
- Word embeddings and visualization
 - https://web.stanford.edu/class/cs224n/assignments/a1_preview/exploring_word_vectors
 - <https://machinelearningmastery.com/develop-word-embeddings-python-gensim/>